LDBC Social Network Benchmark (SNB) – 0.3.0

Coordinator: [names of main authors]

Abstract

LDBC's Social Network Benchmark (LDBC-SNB) is an effort intended to test various functionalities of systems used for graph-like data management. For this, LDBC-SNB uses the recognizable scenario of operating a social network, characterized by its graph-shaped data.

LDBC-SNB consists of three sub-benchmarks, or workloads, that focus on different functionalities. In this document, we present a preliminary version of the Interactive Workload. The other workloads, still in development, are the Business Intelligence Workloads (with analytical queries), and the Graph Analytics Workload (with graph algorithms).

This document contains a detailed explanation of the data used in the whole LDBC-SNB benchmark, a detailed description for all the Interactive Workload, and instructions on how to generate the data and run the benchmark with the provided software.

Executive Summary

The new data economy era, based on complexly structured, distributed and large datasets, has brought on new demands on data management and analytics. As a consequence, new industry actors have appeared, offering technologies specially built for the management of graph-like data. Also, traditional database technologies, such as relational databases, are being adapted to the new demands to remain competitive.

LDBC's Social Network Benchmark (LDBC-SNB) is an industrial and academic initiative, formed by principal actors in the field of graph-like data management. Its goal is to define a framework where different graph based technologies can be fairly tested and compared, that can drive the identification of systems' bottlenecks and required functionalities, and can help researchers to open new research frontiers.

The philosophy around which LDBC-SNB is designed is to be easy to understand, flexible and cheap to adopt. For all these reasons, LDBC-SNB will propose different workloads representing all the usage scenarios of graph-like database technologies, hence, targeting systems of different nature and characteristics. In order increase its adoption by industry and research institutions, LDBC-SNB provides all necessary software, which are designed to be easy to use and deploy at a small cost.

This document contains:

- A detailed specification of the data used in the whole LDBC-SNB benchmark.

- A detailed specification of the workloads.

- A detailed specification of the execution rules of the benchmark.

- A detailed specification of the auditing rules and the full disclosure report's required contents.

TABLE OF CONTENTS

SPECIAL THANKS

Special thanks to all the people that have contributed to the development of this benchmark:

- Renzo Angles (Universidad de Talca)

- Alex Averbuch (Neo Technologies)

- Peter Boncz (Vrije Universiteit Amsterdam)

- Orri Erling (Google)

- Andrey Gubichev (Google)

- Moritz Kaufmann (Technische Universität München)

- Josep Lluís Larriba Pey (Universitat Politècnica de Catalunya)

- Minh-Duc Pham (Altran)

- Marcus Paradies (SAP)

- Arnau Prat (Sparsity Technologies)

- Mirko Spasić (Openlink Software)

- Norbert Martínez (Huawei Technologies)

- Gábor Szárnyas (MTA-BME Lendület Research Group on Cyber-Physical Systems)

LIST OF FIGURES

LIST OF TABLES

DEFINITIONS

DATAGEN: Is the data generator provided by the LDBC-SNB, which is responsible for generating the data needed to run the benchmark.

DBMS: A DataBase Management System.

LDBC-SNB: Linked Data Benchmark Council Social Network Benchmark.

Query Mix: Refers to the ratio between read and update queries of a workload, and the frequency at which they are issued.

SF (Scale Factor): The LDBC-SNB is designed to target systems of different size and scale. The scale factor determines the size of the data used to run the benchmark, measured in Gigabytes.

SUT: The System Under Test is defined to be the database system where the benchmark is executed.

Test Driver: A program provided by the LDBC-SNB, which is responsible for executing the different workloads and gathering the results.

Test Sponsor: The Test Sponsor is the company officially submitting the Result with the FDR and will be charged the filing fee. Although multiple companies may sponsor a Result together, for the purposes of the LDBC processes the Test Sponsor must be a single company. A Test Sponsor need not be a LDBC member. The Test Sponsor is responsible for maintaining the FDR with any necessary updates or corrections. The Test Sponsor is also the name used to identify the Result.

Workload: A workload refers to a set of queries of a given nature (i.e. interactive, analytical, business), how they are issued and at which rate.

1.1 Motivation for the Benchmark

The new era of data economy, based on large, distributed and complexly structured data sets, has brought on new and complex challenges in the field of data management and analytics. These data sets, usually modeled as large graphs, have attracted both industry and academia, due to new opportunities in research and innovation they offer. This situation has also opened the door for new companies to emerge, offering new non-relational and graph-like technologies that are called to play a significant role in upcoming years.

The change in the data paradigm calls for new benchmarks to test new emerging technologies, as they set a framework where different systems can compete and be compared in a fair way, they let technology providers identify the bottlenecks and gaps of their systems and, in general, drive the research and development of new information technology solutions. Without them, the uptake of these technologies is at risk by not providing the industry with clear, user-driven targets for performance and functionality.

The LDBC Social Network Benchmark (LDBC-SNB) aims at being comprehensive benchmark setting the rules for the evaluation of graph-like data management technologies. LDBC-SNB is designed to be a plausible look-alike of all the aspects of operating a social network site, as one of the most representative and relevant use case of modern graph-like applications. LDBC-SNB is a work in progress, and initially, it only includes the Interactive Workload [2], but two more workloads will be introduced in the future: Business Intelligence and Analytics. By designing three separate workloads, LDBC-SNB targets a broader range of systems with different nature and characteristics. LDBC-SNB aims at capturing the essential features of these usage scenarios while abstracting away details of specific business deployments.

1.2 Relevance to Industry

LDBC-SNB is intended to provide the following value to different stakeholders:

- For **end users** facing graph processing tasks, LDBC-SNB provides a recognizable scenario against which it is possible to compare merits of different products and technologies. By covering a wide variety of scales and price points, LDBC-SNB can serve as an aid to technology selection.

- For **vendors** of graph database technology, LDBC-SNB provides a checklist of features and performance characteristics that helps in product positioning and can serve to guide new development.

- For **researchers**, both industrial and academic, the LDBC-SNB dataset and workload provide interesting challenges in multiple choke-point areas, such as query optimization, (distributed) graph analysis, transactional throughput, and provides a way to objectively compare the effectiveness and efficiency of new and existing technology in these areas.

The technological scope of LDBC-SNB comprises all systems that one might conceivably use to perform social network data management tasks:

- **Graph database systems** (e.g. Neo4j, InfiniteGraph, Sparksee, Titan) are novel technologies aimed at storing directed and labeled graphs. They support graph traverals, typically by means of APIs, though some of them also support some sort of graph oriented query language (e.g. Neo4j's Cypher). These systems' internal structures are typically designed to store dynamic graphs that change over time. They offer support transactional queries with some degree of consistency, and value-based indexes to quickly locate nodes and edges. Finally, their architecture is typically single-machine (non-cluster). These systems can potentially implement all three workloads, though Interactive and Business Intelligence workloads are where they will presumably be more competitive.

- **Graph programming frameworks** (e.g. Giraph, Signal/Collect, GraphLab, Green Marl) are designed to perform global graph queries computations, executed in parallel or lockstep. These computations are typically long latency, involving many nodes and edges and often consist of approximation answers to NP-complete problems. These systems expose an API, sometimes following a vertex-centric paradigm, and their architecture targets both single-machine and cluster systems. Though these systems will likely implement the Graph Analytics workload.

- **RDF database systems** (e.g. OWLIM, Virtuoso, BigData, Jena TDB, Stardog, Allegrograph) are systems that implement the SPARQL 1.1 query language, similar in complexity to SQL-92, which allows for structured queries, and simple traversals. RDF database system often come with additional support for simple reasoning (sameAs, subClass), text search and geospatial predicates. RDF database systems generally support transactions, but not always with full concurrency and serializability and their supposed strength is integrating multiple data sources (e.g. DBpedia). Their architecture is both single-machine and clustered, and they will likely target Interactive and Business Intelligence workloads.

- **Relational database systems** (e.g. Postgres, MySQL, Oracle, DB2, SQLserver, Virtuoso, MonetDB, Vectorwise, Vertica, but also Hive and Impala) treat data as relational, and queries are formulated in SQL and/or PL/SQL. Both single-machine and cluster systems exist. They do not normally support recursion, or stateful recursive algorithms, which makes them not at home in the Graph Analytics workloads

- **noSQL database systems** (e.g. key-value stores such as HBase, REDIS, MongoDB, CouchDB, or even MapReduce systems like Hadoop and Pig) are cluster-vbased and scalable. Key-value stores could possibly implement the Interactive Workload, though its navigational aspects would pose some problems as potentially many key-value lookups are needed. MapReduce systems could be suited for the Graph Analytics workload, but their query latency would presumably be so high that the Business Intelligence workload would not make sense, though we note that some of the key-value stores (e.g. MongoDB) provide a MapReduce query functionality on the data that it stores which could make it suited for the Business Intelligence workload.

1.3 General Benchmark Overview

LDBC-SNB aims at being a complete benchmark, designed with the following goals in mind:

- **Rich coverage**. LDBC-SNB is intended to cover most demands encountered in the management of complexly structured data.

- **Modularity**. LDBC-SNB is broken into parts that can be individually addressed. In this manner LDBC-SNB stimulates innovation without imposing an overly high threshold for participation.

- **Reasonable implementation cost**. For a product offering relevant functionality, the effort for obtaining initial results with SNB should be small, on the order of days.

- **Relevant selection of challenges**. Benchmarks are known to direct product development in certain directions. LDBC-SNB is informed by the state of the art in database research so as to offer optimization challenges for years to come while not having a prohibitively high threshold for entry.

- **Reproducibility and documentation of results**. LDBC-SNB will specify the rules for full disclosure of benchmark execution and for auditing of benchmark runs. The workloads may be run on any equipment but the exact configuration and price of the hardware and software must be disclosed.

LDBC-SNB benchmark is modeled around the operation of a real social network site. A social network site represents a relevant use case for the following reasons:

- It is simple to understand for a large audience, as it is arguably present to our every-day life in different shapes and forms.

- It allows testing a complete range of interesting challenges, by means of different workloads targeting systems of different nature and characteristics.

- A social network can be scaled, allowing the design of a scalable benchmark targeting systems of different sizes and budgets.

In Section 2.3, LDBC-SNB defines the schema of the data used in the benchmark. The schema, represents a realistic social network, including people and their activity in the social network during a period of time. Personal information of each person, such as name, birthday, interests or places where people work or study, is included. Persons' activity is represented in the form of friendhisp relationships and content sharing (i.e. messages and pictures). LDBC-SNB provides a scalable synthetic data generator based on the MapReduce parallel paradigm, that produces networks with the described schema with distributions and correlations similar to those expected in a real social network. Furthermore, the data generator is designed to be user-friendly. The proposed data schema is shared by all the different proposed workloads, those we currently have, and those that will be proposed in the future.

In Section 3, the Interactive Workload is proposed. Two more workloads are planned: Business Intelligence Workload and Analytical workload. Workloads are designed to mimic the different usage scenarios found in operating a real social network site, and each of them targets one or more types of systems. Each workload defines a set of queries and query mixes, designed to stress the SUTs in different choke-point areas, while being credible and realistic. Interactive workload reproduces the interaction between the users of the social network by including lookups and transactions that update small portions of the data base. These queries are designed to be interactive and target systems capable of responding such queries with low latency for multiple concurrent users. Business Intelligence workload, will represent those business intelligence analytics a social network company would like to perform in the social network, in order to take advantage of the data to discover new business opportunities. This workload will explore moderate portions of data from different entities, and performing more resource intensive operations. Finally, the graph analytics workload will aim at exploring the characteristics of the underlying structure of the network. Shortest paths, community detection or centrality, are representative queries of this workload, and will imply touching a vast amount of the dataset.

LDBC-SNB provides an execution test driver, which is responsible for executing the workloads and gathering the results. The driver is designed with simplicity and portability in mind, to ease the implementation on systems with different nature and characteristics, at a low implementation cost. Furthermore, it automatically handles the validation of the queries by means of a validation dataset provided by LDBC. The overall philosophy of LDBC-SNB is to provide all the necessary software tools to run the benchmark, and therefore to reduce the benchmark's entry point as much as possible.

1.4   Related Projects

LDBC [1] provides other benchmarks:

- The Semantic Publishing Benchmark (SPB)[1] measures the performance of RDF engines operating on semantic data sets.

- The Graphalytics benchmark [4] measures the performance of graph analysis operations (e.g. PageRank, local clustering coefficient).

---

[1]`http://ldbcouncil.org/benchmarks/spb`

1.5 Participation of Industry and Academia

The list of institutions that take part in the definition and development of LDBC-SNB is formed by relevant actors from both the industry and academia in the field of linked data management. All the participants have contributed with their experience and expertise in the field, making a credible and relevant benchmark that meets all the desired needs. The list of participants is the following:

- FOUNDATION FOR RESEARCH AND TECHNOLOGY HELLAS

- MTA-BME LENDUELET RESEARCH GROUP ON CYBER-PHYSICAL SYSTEMS

- NEO4J

- ONTOTEXT

- OPENLINK

- TECHNISCHE UNIVERSITAET MUENCHEN

- UNIVERSITAET INNSBRUCK

- UNIVERSITAT POLITECNICA DE CATALUNYA

- VRIJE UNIVERSITEIT AMSTERDAM

Besides the aforementioned institutions, during the development of the benchmark several meetings with the technical and users community have been conducted, receiving an invaluable feedback that has contributed to the whole development of the benchmark in every of its aspects.

2.1 Requirements

LDBC-SNB is designed to be flexible and to have an affordable entry point. From small single node and in memory systems to large distributed multi-node clusters have its own place in LDBC-SNB. Therefore, the requirements to fulfill for executing LDBC-SNB are limited to pure software requirements to be able to run the tools. All the software provided by LDBC-SNB have been developed and tested under Linux-based operating systems.

LDBC-SNB does not impose the usage of any specific type of system, as it targets systems of different nature and characteristics, from graph databases, graph processing frameworks and RDF systems, to traditional relational database management systems. Consequently, any language or API capable of expressing the proposed queries can be used. Similarly, data can be stored in the most convenient manner the test sponsor may decide.

2.2 Software and Useful links

- **LDBC Driver 0.3 –** `https://github.com/ldbc/ldbc_driver`: The driver responsible for executing the LDBC-SNB workload.

- **DATAGEN 0.2.5 –** `https://github.com/ldbc/ldbc_snb_datagen`: The data generator used to generate the datasets of the benchmark.

2.3 Data

This section introduces the data used by LDBC-SNB. This includes the different data types, the data schema, how it is generated and the different scale factors.

2.3.1 Data Types

Table 2.1 describes the different types used in the whole benchmark.

Type	Description
ID	integer type with 64-bit precision. All IDs within a single entity are unique
32-bit Integer	integer type with 32-bit precision
64-bit Integer	integer type with 64-bit precision
String	variable length text of size 40 Unicode characters
Long String	variable length text of size 256 Unicode characters
Text	variable length text of size 2000 Unicode characters
Date	date with a precision of a day, encoded as a string with the following format: *yyyy-mm-dd*, where *yyyy* is a four-digit integer representing the year, the year, *mm* is a two-digit integer representing the month and dd is a two-digit integer representing the day.
DateTime	date with a precision of milliseconds, encoded as a string with the following format: *yyyy-mm-ddTHH:MM:ss.sss+0000*, where *yyyy* is a four-digit integer representing the year, the year, *mm* is a two-digit integer representing the month and *dd* is a two-digit integer representing the day, *HH* is a two-digit integer representing the hour, *MM* is a two digit integer representing the minute and *ss.sss* is a five digit fixed point real number representing the seconds up to millisecond precision. Finally, the *+0000* of the end represents the timezone, which in this case is always GMT.

Table 2.1: Description of the data types.

2.3.2 Data Schema

Figure 2.1 shows the data schema in UML. The schema defines the structure of the data used in the benchmark in terms of entities and their relations. Data represents a snapshot of the activity of a social network during a period of time. Data includes entities such as Persons, Organizations, and Places. The schema also models the way persons interact, by means of the friendship relations established with other persons, and the sharing of content such as messages (both textual and images), replies to messages and likes to messages. People form groups to talk about specific topics, which are represented as tags.

LDBC-SNB has been designed to be flexible and to target systems of different nature and characteristics. As such, it does not force any particular internal representation of the schema. The DATAGEN component supports multiple output data formats to fit the needs of different types of systems, including RDF, relational DBMS and graph DBMS.

Figure 2.1: The LDBC-SNB data schema
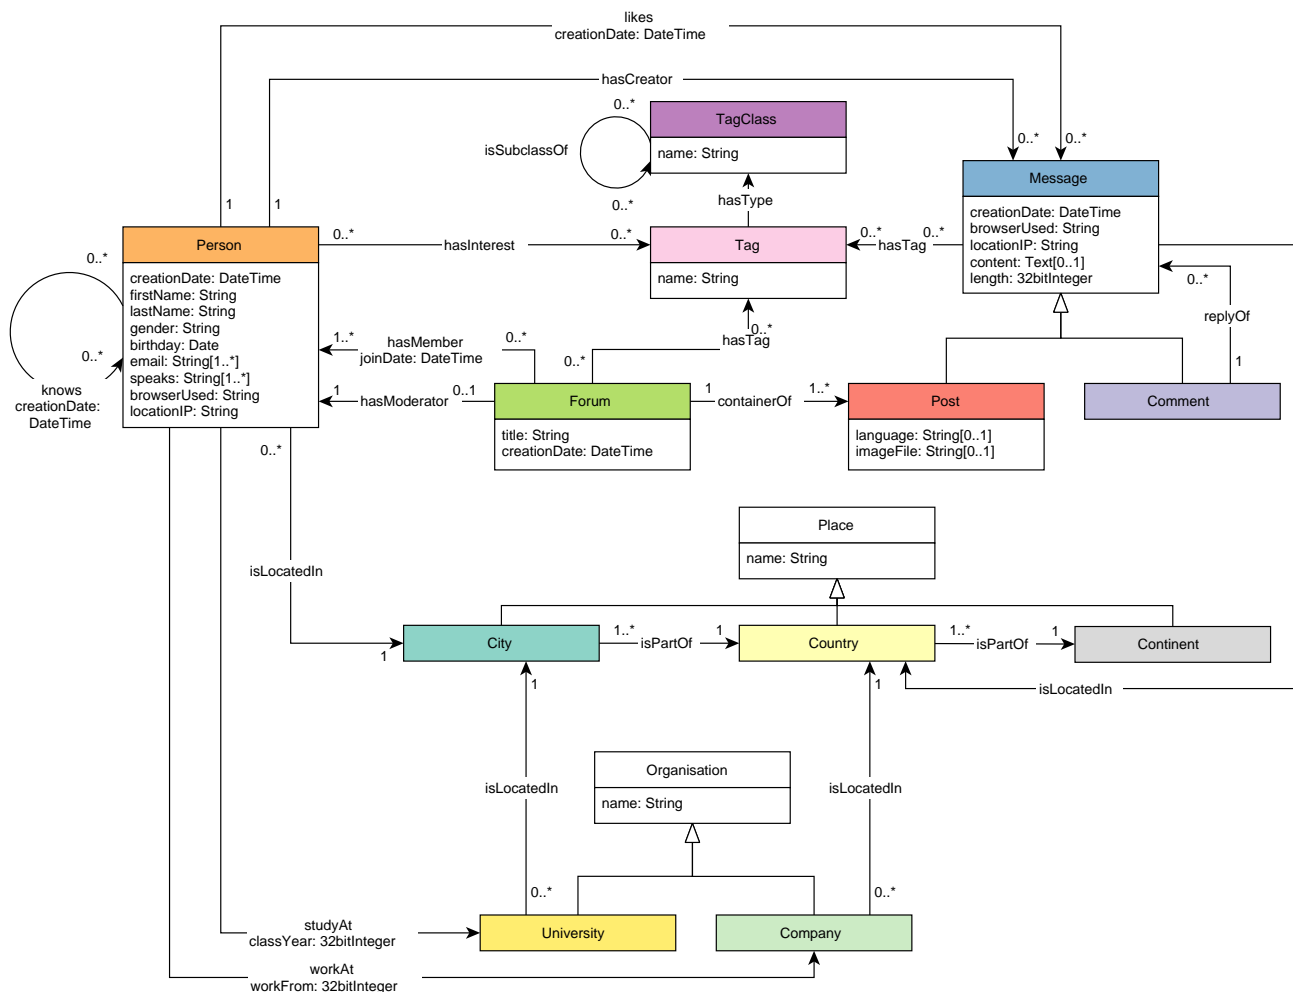
The schema specifies different entities, their attributes and their relations. All of them are described in the following sections.

Textual Restrictions and Notes

- Posts have content or imageFile. They have one of them but not both. The one they do not have is an empty string.

- Posts in a forum can be created by a non-member person if and only if that person is a modeartor.

Entities

City: a sub-class of a Place, and represents a city of the real world. City entities are used to specify where persons live, as well as where universities operate.

Comment: a sub-class of a Message, and represents a comment made by a person to an existing message (either a Post or a Comment).

Company: a sub-class of an Organization, and represents a company where persons work.

Country: a sub-class of a Place, and represents a continent of the real world.

Forum: a meeting point where people post messages. Forums are characterized by the topics (represented as tags) people in the forum are talking about. Although from the schema's perspective it is not evident, there exist three different types of forums: persons' personal walls, image albums, and groups. They are distinguished by their titles. Table 2.2 shows the attributes of Forum entity.

Attribute	Type	Description
id	ID	The identifier of the forum.
title	Long String	The title of the forum.
creationDate	DateTime	The date the forum was created

Table 2.2: Attributes of Forum entity.

Message: an abstract entity that represents a message created by a person. Table 2.3 shows the attributes of Message abstract entity.

Attribute	Type	Description
id	ID	The identifier of the message.
browserUsed	String	The browser used by the Person to create the message.
creationDate	DateTime	The date the message was created.
locationIP	String	The IP of the location from which the message was created.
content	Text[0..1]	The content of the message.
length	32-bit Integer	The length of the content.

Table 2.3: Attributes of Message interface.

Organization: an institution of the real world. Table 2.4 shows the attributes of Organization entity.

Attribute	Type	Description
id	ID	The identifier of the organization.
name	Long String	The name of the organization.

Table 2.4: Attributes of Organization entity.

Person: the avatar a real world person creates when he/she joins the network, and contains various information about the person as well as network related information. Table 2.5 shows the attributes of Person entity.

Attribute	Type	Description
id	ID	The identifier of the person.
firstName	String	The first name of the person.
lastName	String	The last name of the person.
gender	String	The gender of the person.
birthDay	Date	The birthday of the person .
email	Long String[1..*]	The set of emails the person has.
speaks	String[1..*]	The set of languages the person speaks.
browserUser	String	The browser used by the person when he/she registered to the social network.
locationIp	String	The IP of the location from which the person was registered to the social network.
creationDate	DateTime	The date the person joined the social network.

Table 2.5: Attributes of Person entity.

Place: a place in the world. Table 2.6 shows the attributes of Place entity.

Attribute	Type	Description
id	ID	The identifier of the place.
name	Long String	The name of the place.

Table 2.6: Attributes of Place entity.

Post: a sub-class of Message, that is posted in a forum. Posts are created by persons into the forums where they belong. Posts contain either content or imageFile, always one of them but never both. The one they do not have is an empty string. Table 2.7 shows the attributes of Post entity.

Attribute	Type	Description
language	String[0..1]	The language of the post.
imageFile	String[0..1]	The image file of the post..

Table 2.7: Attributes of Post entity.

Tag: a topic or a concept. Tags are used to specify the topics of forums and posts, as well as the topics a person is interested in. Table 2.8 shows the attributes of Tag entity.

Attribute	Type	Description
id	ID	The identifier of the tag.
name	Long String	The name of the tag.

Table 2.8: Attributes of Tag entity.

TagClass: a class or a category used to build a hierarchy of tags. Table 2.9 shows the attributes of TagClass entity.

Attribute	Type	Description
id	ID	The identifier of the tagclass.
name	Long String	The name of the tagclass.

Table 2.9: Attributes of TagClass entity.

University: a sub-class of Organization, and represents an institution where persons study.

Relations

Relations connect entities of different types. Entities are defined by their "id" attribute.

Name	Tail	Head	Type	Description
containerOf	Forum[1]	Post[1..*]	D	A Forum and a Post contained in it
hasCreator	Message[0..*]	Person[1]	D	A Message and its creator (Person)
hasInterest	Person[0..*]	Tag[0..*]	D	A Person and a Tag representing a topic the person is interested in
hasMember	Forum[0..*]	Person[1..*]	D	A Forum and a member (Person) of the forum • **Attribute**: joinDate • **Type:** DateTime • **Description:** The Date the person joined the forum
hasModerator	Forum[0..*]	Person[1]	D	A Forum and its moderator (Person)
hasTag	Message[0..*]	Tag[0..*]	D	A Message and a Tag representing the message's topic
hasTag	Forum[0..*]	Tag[0..*]	D	A Forum and a Tag representing the forum's topic
hasType	Tag[0..*]	TagClass[0..*]	D	A Tag and a TagClass the tag belongs to
isLocatedIn	Company[0..*]	Country[1]	D	A Company and its home Country
isLocatedIn	Message[0..*]	Country[1]	D	A Message and the Country from which it was issued
isLocatedIn	Person[0..*]	City[1]	D	A Person and their home City
isLocatedIn	University[0..*]	City[1]	D	A University and the City where the university is
isPartOf	City[1..*]	Country[1]	D	A City and the Country it is part of
isPartOf	Country[1..*]	Continent[1]	D	A Country and the Continent it is part of
isSubclassOf	TagClass[0..*]	TagClass[0..*]	D	A TagClass its parent TagClass
knows	Person[0..*]	Person[0..*]	U	Two Persons that know each other • **Attribute**: creationDate • **Type:** DateTime • **Description:** The date the knows relation was established
likes	Person[0..*]	Message[0..*]	D	A Person that likes a Message • **Attribute**: creationDate • **Type:** DateTime • **Description:** The date the like was issued
replyOf	Comment[0..*]	Message[1]	D	A Comment and the Message it replies

studyAt	Person[0..*]	University[0..*]	D	A Person and a University it has studied • **Attribute**: classYear • **Type:** 32-bit Integer • **Description:** The year the person graduated.
workAt	Person[0..*]	Company[0..*]	D	A Person and a Company it works • **Attribute:** workFrom • **Type:** 32-bit Integer • **Description:** The year the person started to work at that company

Table 2.10: Description of the data relations.

2.3.3 Output Data

DATAGEN produces outputs three different items:

- **Dataset**: The dataset to be bulk loaded by the SUT. It corresponds to roughly the 90% of the total generated network.

- **Update Streams**: A set of update streams containing update queries, which are used by the driver to generate the update queries of the workloads. This update streams correspond to the remaining 10% of the generated dataset.

- **Substitution Parameters**: A set of files containing the different parameter bindings that will be used by the driver to generate the read queries of the workloads.

The SUT have to take care only of the generated Dataset to be bulk loaded. Three different formats are supported by DATAGEN:

- **CSV:** Data output in CSV format, one file per different entity and on file per different relation. Also, there is a file por those attributes whose cardinality is larger than one (i.e. Person.email, Person.speaks, etc.).

- **CSVMergeForeign:** Similar to CSV format, but in this case, those relations of the form 1 to 1 and 1 to N, are stored in the tail entity file as a foreign keys.

- **Turtle:** Dataset in Turtle format for RDF systems.

CSV

This is a comma separated format. Each entity, relation and properties with a cardinality larger than one, are output in a separate file. Generated files are summarized at Table 2.11. Depending on the number of threads used for generating the dataset, the number of files varies, since there is a file generated per thread. The * in the file names indicates a number between 0 and $NumberOfThreads - 1$.

File	Content
comment_*.csv	id | creationDate | locationIP | browserUsed | content | length
comment_hasCreator_person_*.csv	Comment.id | Person.id
comment_isLocatedIn_place_*.csv	Comment.id | Place.id
comment_replyOf_comment_*.csv	Comment.id | Comment.id
comment_replyOf_post_*.csv	Comment.id | Post.id
forum_*.csv	id | title | creationDate
forum_containerOf_post_*.csv	Forum.id | Post.id
forum_hasMember_person_*.csv	Forum.id | Person.id | joinDate
forum_hasModerator_person_*.csv	Forum.id | Person.id
forum_hasTag_tag_*.csv	Forum.id | Tag.id
organization_*.csv	id | type("university", "company") | name | url
organisation_isLocatedIn_place_*.csv	Organisation.id | Place.id
person_*.csv	id | firstName | lastName | gender | birthday | creationDate | locationIP | browserUsed
person_email_emailaddress_*.csv	Person.id | email
person_hasInterest_tag_*.csv	Person.id | Tag.id
person_isLocatedIn_place_*.csv	Person.id | Place.id
person_knows_person_*.csv	Person.id | Person.id | creationDate
person_likes_comment_*.csv	Person.id | Post.id | creationDate
person_likes_post_*.csv	Person.id | Post.id | creationDate
person_speaks_language_*.csv	Person.id | language
person_studyAt_organisation_*.csv	Person.id | Organisation.id | classYear
person_workAt_organisation_*.csv	Person.id | Organisation.id | workFrom
place_*.csv	id | name | url | type("city", "country", "continent")
place_isPartOf_place_*.csv	Place.id | Place.id
post_*.csv	id | imageFile | creationDate | locationIP | browserUsed | language | content | length
post_hasCreator_person_*.csv	Post.id | Person.id
post_hasTag_tag_*.csv	Post.id | Tag.id
post_isLocatedIn_place.csv	Post.id | Place.id
tag_*.csv	id | name | url
tag_hasType_tagclass_*.csv	Tag.id | TagClass.id
tagclass_*.csv	id | name | url
tagclass_isSubclassOf_tagclass_*.csv	TagClass.id | TagClass.id

Table 2.11: Files output by the CSV serializer (32 in total)

CSV_MERGE_FOREIGN

This is a comma separated format. It is similar to CSV, but those relations connecting two entities A and B, where an entity A has a cardinality of one, A is output as a column of entity B. Generated files are summarized at Table 2.12. Depending on the number of threads used for generating the dataset, the number of files varies, since there is a file generated per thread. The * in the file names indicates a number between 0 and $NumberOfThreads - 1$.

Turtle

This is the standard Turtle[1] format. DATAGEN outputs two files: `0_ldbc_socialnet_static_dbp.ttl` and `0_ldbc_socialnet.ttl`.

2.3.4   Scale Factors

LDBC-SNB defines a set of scale factors (SFs), targeting systems of different sizes and budgets. SFs are computed based on the ASCII size in Gigabytes of the generated output files using the CSV serializer. For example, SF 1 weights roughly 1 GB in CSV format, SF 3 weights roughly 3 GB and so on and so forth. The proposed SFs are the following: 1, 3, 10, 30, 100, 300, 1000. The Test Sponsor may select the SF that better fits their needs, by properly configuring the DATAGEN.

The size of the resulting dataset, is mainly affected by the following configuration parameters: the number of persons and the number of years simulated. Different SFs are computed by scaling the number of Persons in the network, while fixing the number of years simulated. Table 2.13 shows the parameters used in each of the SFs.

---

[1]Description of the Turtle RDF format http://www.w3.org/TR/turtle/

File	Content
comment_*.csv	id | creationDate | locationIP | browserUsed | content | length | creator | place | replyOfPost | replyOfComment
forum_*.csv	id | title | creationDate | moderator
forum_hasMember_person_*.csv	Forum.id | Person.id | joinDate
forum_hasTag_tag_*.csv	Forum.id | Tag.id
organization_*.csv	id | type("university", "company") | name | url
organisation_isLocatedIn_place_*.csv	Organisation.id | Place.id
person_*.csv	id | firstName | lastName | gender | birthday | creationDate | locationIP | browserUsed | place
person_email_emailaddress_*.csv	Person.id | email
person_hasInterest_tag_*.csv	Person.id | Tag.id
person_knows_person_*.csv	Person.id | Person.id | creationDate
person_likes_comment_*.csv	Person.id | Post.id | creationDate
person_likes_post_*.csv	Person.id | Post.id | creationDate
person_speaks_language_*.csv	Person.id | language
person_studyAt_organisation_*.csv	Person.id | Organisation.id | classYear
person_workAt_organisation_*.csv	Person.id | Organisation.id | workFrom
place_*.csv	id | name | url | type("city", "country", "continent")
place_isPartOf_place_*.csv	Place.id | Place.id
post_*.csv	id | imageFile | creationDate | locationIP | browserUsed | language | content | length | creator | Forum.id | place
post_hasTag_tag_*.csv	Post.id | Tag.id
tag_*.csv	id | name | url
tag_hasType_tagclass_*.csv	Tag.id | TagClass.id
tagclass_*.csv	id | name | url
tagclass_isSubclassOf_tagclass_*.csv	TagClass.id | TagClass.id

Table 2.12: Files output by the CSV_MERGE_FOREIGN serializer (23 in total)

Scale Factor	1	3	10	30	100	300	1000
# of Persons	11K	27K	73K	182K	499K	1.25M	3.6M
# of Years	3	3	3	3	3	3	3
Start Year	2010	2010	2010	2010	2010	2010	2010

Table 2.13: Parameters of each scale factor.

For example, SF 30 consists of the activity of a social network of 182K users during a period of three years, starting from 2010.

3.1 Query Description Format

Queries are described in natural language using a well-defined structure that consists of three sections: *description*, a concise textual description of the query; *parameters*, a list of input parameters and their types; and *results*, a list of expected results and their types. The syntax used in *parameters* and *results* sections is as follows:

- **Entity**: entity type in the dataset.
  One word, possibly constructed by appending multiple words together, starting with uppercase character and following the camel case notation, e.g. `TagClass` represents an entity of type "TagClass".

- **Relationship**: relationship type in the dataset.
  One word, possibly constructed by appending multiple words together, starting with lowercase character and following the camel case notation, and surrounded by arrow to communicate direction, e.g. `-worksAt->` represents a directed relationship of type "worksAt".

- **Attribute**: attribute of an entity or relationship in the dataset.
  One word, possibly constructed by appending multiple words together, starting with lowercase character and following the camel case notation, and prefixed by a "." to dereference the entity/relationship, e.g. `Person.firstName` refers to "firstName" attribute on the "Person" entity, and `-studyAt->.classYear` refers to "classYear" attribute on the "studyAt" relationship.

- **Unordered Set**: an unordered collection of distinct elements.
  Surrounded by { and } braces, with the element type between them, e.g. `{String}` refers to a set of strings.

- **Ordered List**: an ordered collection where duplicate elements are allowed.
  Surrounded by [and] braces, with the element type between them, e.g. `[String]` refers to a list of strings.

- **Ordered Tuple**: a fixed length, fixed order list of elements, where elements at each position of the tuple have predefined, possibly different, types.
  Surrounded by < and > braces, with the element types between them in a specific order e.g. `<String, Boolean>` refers to a 2-tuple containing a string value in the first element and a boolean value in the second, and `[<String, Boolean>]` is an ordered list of those 2-tuples.

3.2 Graph Patterns

We use the following notation for defining graph patterns:

- Filtering conditions are typeset in *italic*.

- Properties that should be returned are denoted in normal (book) font.

- Negative conditions, i.e., edges that are now allowed in the graph are denoted with dashed red lines.
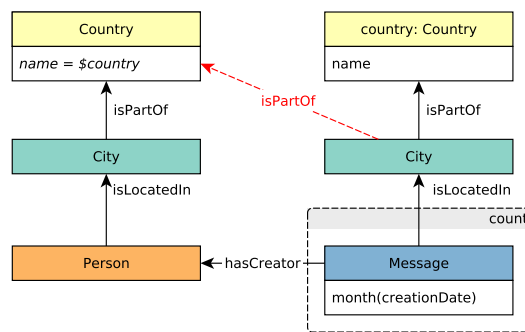
- Aggregations are shown in dashed boxes.

Figure 3.1: Example graph pattern.

4.1 Choke Points

The design of the interactive workload queries has been conceived around two main aspects: realism and technological relevance. While realism has been assessed by looking at existing social networks and thinking about what interesting functionalities a user might desire from them, technological relevance has been achieved by identifying a set of choke points queries should stress. These choke points capture those critical operations, techniques or technologies that could significatively affect the performance of the queries. The choke points can be summarized in the following list:

4.1.1 Aggregation Performance

The queries generally have a top-k order by and often a group by in addition to this. These offer multiple optimization opportunities. The queries also often have distinct operators, i.e. distinct friends within two steps. Collectively these are all set operations that may be implemented with some combination of hash and sorting, possibly exploiting ordering in the data itself. The aggregates are not limited to counts and sums. For example string concatenation occurs as an aggregate, testing possible user defined aggregate support. There is a wide range of cardinalities in grouping, from low, e.g. country, to high, e.g. post.

4.1.2 Join Performance

Each graph traversal step is in principle a join. The join patterns are diverse, exercising both index- and hash-based operators. Queries are designed so as to reward judicious use of hash join by having patterns starting with one entity, fetching many related entities and then testing how these are related to a third entity, e.g. posts of a user with a tag of a given type.

4.1.3 Data Access Locality

Graph problems are notoriously non-local. However, when queries touch any non-trivial fraction of a dataset, locality will emerge and can be exploited, for example by vectored index access or by ordering data so that that a merge join is possible.

4.1.4 Expression Calculation

Queries often have expressions, including conditional expressions. This provides opportunities for vectoring and tests efficient management of intermediate results.

4.1.5 Correlated Subqueries

The workload has many correlated subqueries, for example constructts like x within two steps but not in one step, which would typically be a correlated subquery with NOT EXISTS. There are also scalar subqueries with aggregation, for example returning the count of posts satisfying a certain criteria.

4.1.6 Parallelism and Concurrency

All queries offer opportunities for parallelism. This tests a wide range of constructs, for example partitioned parallel variants of group by and distinct. An interactive workload will typically avoid trivially parallelizable table scans. Thus the opportunities that exist must be captured by index based, navigational query plans. The choice of whether to parallelize or not is often left to run time and will have to depend on the actual data seen in the execution, as starting a parallel thread with too little to do is counter-productive.

4.1.7 Graph Specifics

Graph problems are generally characterized by transitive properties and the fact that neighboring vertices often have a large overlap in their environments. This makes cardinality estimation harder. For example, a query optimizer needs to recognize whether a relationship has a tree or graph shape in order to make correct cardinality estimations. Further, there are problems aggregating properties over a set of consecutive edges. The workload contains business questions dealing with paths and aggregates across paths, as well as the easier case of determining a membership in a hierarchy with a transitive part-of relation.

4.2 Query Specifications

4.2.1 Complex Reads Query Descriptions

Notes:

- Some queries require returning the content of a post. As stated in the schema, posts have either some content or an imageFile, but not both. An empty string in content represents the post not having any content, therefore, it must have a non-empty string in imageFile or the other way around.

1. Friends with certain name

   - **Description:** Given a start Person, find Persons with a given first name that the start Person is connected to (excluding start Person) by at most 3 steps via Knows relationships. Return Persons, including summaries of the Persons workplaces and places of study.
   - **Parameters:**

Person.id	ID
Person.firstName	String

   - **Results:**

Person.id	ID
Person.lastName	String
Person.birthday	Date
Person.creationDate	DateTime
Person.gender	String
Person.browserUsed	String
Person.locationIP	String
{Person.emails}	{String}
{Person.language}	{String}
Person-isLocatedIn->Place.name	String
{Person-studyAt->University.name, Person-studyAt->.classYear, Person-studyAt->University-isLocatedIn->City.name}	{<String, 32-bit Integer, String>}
{Person-workAt->Company.name, Person-workAt->.workFrom, Person-workAt->Company-isLocatedIn->Country.name}	{<String, 32-bit Integer, String>}

   - **Sort:**

1st	distance from person (ascending)
2nd	Person.lastName (ascending)
3rd	Person.id (ascending)

   - **Limit:**

     20

2. Recent posts and comments by your friends

   - **Description:** Given a start Person, find (most recent) Messages from all of that Person's friends, that were created before (and including) a given date.

- **Parameters:**
Person.id	ID
date	DateTime

- **Results:**
Message-hasCreator->Person.id	ID
ssage-hasCreator->Person.firstName	String
ssage-hasCreator->Person.lastName	String
Message.id	ID
Message.content or Post.imageFile	String
Message.creationDate	DateTime

- **Sort:**

  1st Message.creationDate (descending)

  2nd Message.id (ascending)

- **Limit:**

  20

3. Friends and friends of friends that have been to countries X and Y

- **Description:** Given a start Person, find Persons that are their friends and friends of friends (excluding start Person) that have made Posts/Comments in both of the given Countries, X and Y, within a given period. Only Persons that are foreign to Countries X and Y are considered, that is Persons whose Location is not Country X or Country Y.

- **Parameters:**
Person.id	ID	
CountryX.name	String	
CountryY.name	String	
startDate	Date	// beginning of requested period
duration	32-bit Integer	// duration of requested period, in days
		// the interval [startDate, startDate + Duration) is closed-open

- **Results:**
Person.id	ID	
Person.firstName	String	
Person.lastName	String	
countx	32-bit Integer	// number of Messages from Country X made by Person within the given time
county	32-bit Integer	// number of Messages from Country Y made by Person within the given time
count	32-bit Integer	// countx + county

- **Sort:**

  1st countx (descending)

  2nd Person.id (ascending)

- **Limit:**

  20

4. New topics

- **Description:** Given a start Person, find Tags that are attached to Posts that were created by that Person's friends. Only include Tags that were attached to friends' Posts created within a given time interval, and that were never attached to friends' Posts created before this interval.

- **Parameters:**
Person.id	ID	
startDate	Date	
duration	32-bit Integer	// duration of requested period, in days
		// the interval [startDate, startDate + Duration) is closed-open

- **Results:**
Tag.name	String	
count	32-bit Integer	// number of Posts made within the given time interval that have this Tag

- **Sort:**

  1st count (descending)

  2nd Tag.name (ascending)

- **Limit:**

     10

5. New groups

   - **Description:** Given a start Person, find the Forums which that Person's friends and friends of friends (excluding start Person) became Members of after a given date. For each forum find the number of Posts that were created by any of these Persons. For each Forum and consider only those Persons which joined that particular Forum after the given date.

   - **Parameters:**
Person.id	ID
date	Date

   - **Results:**
Forum.title	String	
count	32-bit Integer	// number of Posts made in Forum that were created by friends

   - **Sort:**

     1st count (descending)

     2nd Forum.id (ascending)

   - **Limit:**

      20

6. Tag co-occurrence

   - **Description:** Given a start Person and some Tag, find the other Tags that occur together with this Tag on Posts that were created by start Person's friends and friends of friends (excluding start Person). Return For each Tag, find the count of Posts that were created by these Persons, which contain both this Tag and the given Tag.

   - **Parameters:**
Person.id	ID
Tag.name	String

   - **Results:**
Tag.name	String	
count	32-bit Integer	// number of Posts that were created by friends and friends of friends, which contain this Tag

   - **Sort:**

     1st count (descending)

     2nd Tag.name (ascending)

   - **Limit:**

      10

7. Recent likers

   - **Description:** Given a start Person, find (most recent) Likes on any of start Person's Messages. Find Persons that Liked any of start Person's Messages, the Messages they liked most recently, creation date of that Like, and the latency (in minutes) between creation of Messages and Like. Additionally, for each Person found return a flag indicating whether the liker is a friend of start Person. In the case that a Person Liked multiple Messages at the same time, return the Message with lowest identifier.

   - **Parameters:**
Person.id	64-bit Integer

- **Results:**

Person.id	ID
Person.firstName	String
Person.lastName	String
Like.creationDate	DateTime
Message.id	ID
Message.content or Post.imageFile	String
latency	32-bit Integer // duration between creation of Message and Like, in minutes
isNew	Boolean // false if liker Person is friend of start Person, true otherwise

- **Sort:**

  1st Like.creationDate (descending)

  2nd Person.id (ascending)

- **Limit:**

  20

8. Recent replies

   - **Description:** Given a start Person, find (most recent) Comments that are replies to Messages of the start Person. Only consider immediate (1-hop) replies, not the transitive (multi-hop) case. Return the reply Comments, and the Person that created each reply Comment.

   - **Parameters:**

Person.id	ID

   - **Results:**

Person.id	ID
Person.firstName	String
Person.lastName	String
Comment.creationDate	DateTime
Comment.id	ID
Comment.content	String

   - **Sort:**

     1st Comment.creationDate (descending)

     2nd Comment.id (ascending)

   - **Limit:**

     20

9. Recent posts and comments by friends or friends of friends

   - **Description:** Given a start Person, find the (most recent) Messages created by that Person's friends or friends of friends (excluding start Person). Only consider the Messages created before a given date (excluding that date).

   - **Parameters:**

Person.id	ID
date	Date

   - **Results:**

Message-hasCreator->Person.id	ID
Message-hasCreator->Person.firstName	String
Message-hasCreatr->Person.lastName	String
Message.id	ID
Message.content or Post.imageFile	String
Message.creationDate	DateTime

   - **Sort:**

     1st Message.creationDate (descending)

     2nd Message.id (ascending)

LDBC Social Network Benchmark (SNB) – 0.3.0

- **Limit:**

  20

10. Friend recommendation

    - **Description:** Given a start Person, find that Person's friends of friends (excluding start Person, and immediate friends), who were born on or after the 21st of a given month (in any year) and before the 22nd of the following month. Calculate the similarity between each of these Persons and start Person, where similarity for any Person is defined as follows:

        – common = number of Posts created by that Person, such that the Post has a Tag that start Person is Interested in

        – uncommon = number of Posts created by that Person, such that the Post has no Tag that start Person is Interested in

        – similarity = common - uncommon

    - **Parameters:**

Person.id	ID
month	32-bit Integer // between 1-12

    - **Results:**

Person.id	ID
Person.firstName	String
Person.lastName	String
similarity	32-bit Integer
Person.gender	String
Person-isLocatedIn->Place.name	Sting

    - **Sort:**

      1st similarity (descending)

      2nd Person.id (ascending)

    - **Limit:**

      10

11. Job referral

    - **Description:** Given a start Person, find that Person's friends and friends of friends (excluding start Person) who started Working in some Company in a given Country, before a given date (year).

    - **Parameters:**

Person.id	ID
Country.name	String
year	32-bit Integer

    - **Results:**

Person.id	ID
Person.firstName	String
Person.lastName	String
Person-worksAt->Organization.name	String
Person-worksAt->.worksFrom	32-bit Integer

    - **Sort:**

      1st Person-worksAt->.worksFrom (ascending)

      2nd Person.id (ascending)

      3st Person-worksAt->Organization.name (descending)

    - **Limit:**

      10

12. Expert search

- **Description:** Given a start Person, find the Comments that this Person's friends made in reply to Posts, considering only those Comments that are immediate (1-hop) replies to Posts, not the transitive (multi-hop) case. Only consider Posts with a Tag in a given TagClass or in a descendent of that TagClass. Count the number of these reply Comments, and collect the Tags (with valid tag class) that were attached to the Posts they replied to. Return Persons with at least one reply, the reply count, and the collection of Tags.
- **Parameters:**

  Person.id   ID

  TagClass.name   String

- **Results:**

  Person.id   ID

  Person.firstName   String

  Person.lastName   String

  {Tag.name}   {String}

  count   32-bit Integer   // number of reply Comments

- **Sort:**

  1st   count (descending)

  2nd   Person.id (ascending)

- **Limit:**

  20

13. Single shortest path

    - **Description:** Given two Persons, find the shortest path between these two Persons in the subgraph induced by the Knows relationships. Return the length of this path.
      - -1 : no path found
      - 0: start person = end person
      - > 0: regular case
    - **Parameters:**

      Person.id   ID   // person 1

      Person.id   ID   // person 2

    - **Results:**

      length   32-bit Integer

    - **Sort:**

      -

    - **Limit:**

      -

14. Weighted/unweighted paths

    - **Description:** Given two Persons, find all (unweighted) shortest paths between these two Persons, in the subgraph induced by the Knows relationship. Then, for each path calculate a weight. The nodes in the path are Persons, and the weight of a path is the sum of weights between every pair of consecutive Person nodes in the path. The weight for a pair of Persons is calculated such that every reply (by one of the Persons) to a Post (by the other Person) contributes 1.0, and every reply (by ones of the Persons) to a Comment (by the other Person) contributes 0.5. Return all the paths with shortest length, and their weights.
    - **Parameters:**

      Person.id   ID   // person 1

      Person.id   ID   // person 2

    - **Results:**

      [Person.id]   [ID]   // Identifiers representing an ordered sequence of the Persons in the path

      weight   64-bit Float

    - **Sort:**

      1st   weight (descending) // The order of paths with the same weight is unspecified

    - **Limit:**

      -

4.2.2 Short Reads Query Descriptions

1. Person Profile

   - **Description:** Given a start Person, retrieve their first name, last name, birthday, IP address, browser, and city of residence.

   - **Parameters:**
     Person.id   ID

   - **Results:**

Person.firstName	String
Person.lastName	String
Person.birthDay	Date
Person.locationIP	String
Person.browserUsed	String
Person-isLocatedIn->Place.id	32-bit Integer
Person.gender	String
Person.creationDate	DateTime

   - **Sort:**

     -

   - **Limit:**

     -

2. Person Recent Messages

   - **Description:** Given a start Person, retrieve the last 10 Messages created by that user. For each message, return that message, the original post in its conversation, and the author of that post. If any of the Messages is a Post, then the original Post will be the same Message, i.e. that Message will appear twice in that result.

   - **Parameters:**
     Person.id   ID

   - **Results:**

Message.id	64-bit Integer
Message.content or Post.imageFile	String
Message.creationDate	DateTime
Post.id or Comment-replyOf*->Post.id	ID
Post-hasCreator->Person.id or Comment-replyOf*->Post-hasCreator->Person.id	ID
Post-hasCreator->Person.firstName or Comment-replyOf*->Post-hasCreator->Person.firstName	String
Post-hasCreator->Person.lastName or Comment-replyOf*->Post-hasCreator->Person.lastName	String

   - **Sort:**

     1st Message.creationDate (descending)

     2nd Message.id (descending)

   - **Limit:**

     -

3. Person Friends

   - **Description:** Given a start Person, retrieve all of their friends, and the date at which they became friends.

   - **Parameters:**
     Person.id   ID

   - **Results:**

Person.id	ID
Person.firstName	String
Person.lastName	String
Knows.creationDate	DateTime

   - **Sort:**

     1st Knows.creationDate (descending)

     2nd Person.id (ascending)

- **Limit:**

    -

4. Message Content

    - **Description:** Given a Message, retrieve its content and creation date.

    - **Parameters:**
      Message.id   ID

    - **Results:**
      Message.creationDate   ID
      Message.content or Post.imageFile   String

    - **Sort:**

        -

    - **Limit:**

        -

5. Message Creator

    - **Description:** Given a Message, retrieve its author.

    - **Parameters:**
      Message.id   ID

    - **Results:**
      Message-hasCreator->Person.id   ID
      Message-hasCreator->Person.firstName   String
      Message-hasCreator->Person.lastName   String

    - **Sort:**

        -

    - **Limit:**

        -

6. Message Forum

    - **Description:** Given a Message, retrieve the Forum that contains it and the Person that moderates that forum. Since comments are not directly contained in forums, for comments, return the forum containing the original post in the thread which the comment is replying to.

    - **Parameters:**
      Message.id   ID

    - **Results:**
      Message<-containerOf-Forum.id   ID
      Message<-containerOf-Forum.title   String
      Message<-containerOf-Forum-hasModerator->Person.id   ID
      Message<-containerOf-Forum-hasModerator->Person.firstName   String
      Message<-containerOf-Forum-hasModerator->Person.lastName   String

    - **Sort:**

        -

    - **Limit:**

        -

7. Message Replies

    - **Description:** Given a Message, retrieve the (1-hop) Comments that reply to it. In addition, return a boolean flag indicating if the author of the reply knows the author of the original message. If author is same as original author, return false for "knows" flag.

    - **Parameters:**
      Message.id   ID

- **Results:**

Message<-replyOf-Comment.id	ID
Message<-replyOf-Comment.content	String
Message<-replyOf-Comment.creationDate	DateTime
Message-hasCreator->Person.id	ID
Message-hasCreator->Person.firstName	String
Message-hasCreator->Person.lastName	String

- **Sort:**

  1st   Message<-replyOf-Comment.creationDate (descending)

  2nd   Message-hasCreator->Person.id (ascending)

- **Limit:**

  -

4.2.3 Update Query Descriptions

1. Add Person

   - **Description:** Add a Person to the social network.
   - **Parameters:**

Person.id	ID
Person.firstName	String
Person.lastName	String
Person.gender	String
Person.birthDay	Date
Person.creationDate	DateTime
Person.locationIp	String
Person.browserUsed	String
Person-isLocatedIn->City.id	ID
Person.speaks	{ String }
Person.emails	{ String }
Person-hasInterest->Tag.id	{ ID }
{ Person-studyAt->University.id, Person-studyAt->.classYear }	{ID, 32-bit Integer}
{ Person-workAt->Company.id, Person-workAt->.workFrom }	{ID, 32-bit Integer}

2. Add Post Like

   - **Description:** Add a Like to a Post of the social network.
   - **Parameters:**

Person.id	ID
Post.id	ID
Person-likes->.creationDate	DateTime

3. Add Comment Like

   - **Description: Add a Like to a Comment of the social network.**
   - **Parameters:**

Person.id	ID
Comment.id	ID
Person-likes->.creationDate	DateTime

4. Add Forum

   - **Description:** Add a Forum to the social network.
   - **Parameters:**

Forum.id	ID	// person 1
Forum.title	String	// person 2
Forum.creationDate	DateTime	
Forum-hasModerator->Person.id	{ ID }	
Forum-hasTag->Tag.id	{ ID }	

5. Add Forum Membership

   - **Description:** Add a Forum membership to the social network.

   - **Parameters:**
Person.id	ID
Person-hasMember->Forum.id	ID
Person-hasMember->.joinDate	DateTime

6. Add Post

   - **Description:** Add a Post to the social network.

   - **Parameters:**
Post.id	ID
Post.imageFile	String
Post.creationDate	DateTime
Post.locationIp	String
Post.browserUsed	String
Post.language	String
Post.content	Text
Post.length	32-bit Integer
Post-hasCreator->Person.id	ID
Forum-containerOf->Post.id	ID
Post-isLocatedIn->Country.id	ID
{Post-hasTag->Tag.id}	{ID}

7. Add Comment

   - **Description:** Add a Comment replying to a Post/Comment to the social network.

   - **Parameters:**
Comment.id	ID	
Comment.creationDate	DateTime	
Comment.locationIp	String	
Comment.browserUsed	String	
Comment.content	Text	
Comment.length	32-bit Integer	
Comment-hasCreator->Person.id	ID	
Comment-isLocatedIn->Country.id	ID	
Comment-replyOf->Post.id	ID	// -1 if the comment is a reply of a comment.
Comment-replyOf->Comment.id	ID	// -1 if the comment is a reply of a post.
{Comment-hasTag->Tag.id}	{ID}	

8. Add Friendship

   - **Description:** Add a friendship relation to the social network

   - **Parameters:**
Person.id	ID	// person 1
Person.id	ID	// person 2
Person-knows->.creationDate	DateTime	

number	1
title	Friends with certain name
	message: Message *creationDate < $date* length year(creationDate)
description	Given a start Person, find Persons with a given first name that the start Person is connected to (excluding start Person) by at most 3 steps via Knows relationships. Return Persons, including summaries of the Persons workplaces and places of study.
parameters	`Person.id` ID `Person.firstName` String

result	`Person.id` ID
	`Person.lastName` String
	`Person.birthday` Date
	`Person.creationDate` DateTime
	`Person.gender` String
	`Person.browserUsed` String
	`Person.locationIP` String
	`{Person.emails}` {String}
	`{Person.language}` {String}
	`Person-isLocatedIn->Place.name` String
	`{Person-studyAt->University.name, Person-studyAt->.classYear, Person-studyAt->Unive`
	`{Person-workAt->Company.name, Person-workAt->.workFrom, Person-workAt->Company-isLo`

sort	`distance from person` ↑ `Person.lastName` ↑ `Person.id` ↑
limit	20
choke points	

4.3 Substitution parameters

Together with the dataset, DATAGEN produces a set of parameters per query type. Parameter generation is designed in such a way that for each query type, all of the generated parameters yield similar runtime behaviour of that query.

Specifically, the selection of parameters for a query template guarantees the following properties of the resulting queries:

P1: the query runtime has a bounded variance: the average runtime corresponds to the behavior of the majority of the queries

P2: the runtime distribution is stable: different samples of (e.g. 10) parameter bindings used in different query streams result in an identical runtime distribution across streams

P3: the optimal logical plan (optimal operator order) of the queries is the same: this ensures that a specific query template tests the system's behavior under the well-chosen technical difficulty (e.g. handling voluminous joins or proper cardinality estimation for subqueries, etc.)

As a result, the amount of data that the query touches is roughly the same for every parameter binding, assuming that the query optimizer figures out a reasonable execution plan for the query. This is done to avoid bindings that cause unexpectedly long or short runtimes of queries, or even result in a completely different optimal execution plan. Such effects could arise due to the data skew and correlations between values in the generated dataset.

In order to get the parameter bindings for each of the queries, we have designed a *Parameter Curation* procedure that works in two stages:

1. for each query template for all possible parameter bindings, we determine the size of intermediate results in the *intended* query plan. Intermediate result size heavily influences the runtime of a query, so two queries with the same operator tree and similar intermediate result sizes at every level of this operator tree are expected to have similar runtimes. This analysis is effectively a side effect of data generation, that is we keep all the necessary counts (number of friends per user, number of posts of friends etc.) as we create the dataset.

2. then, a greedy algorithm selects ("curates") those parameters with similar intermediate result counts from the domain of all the parameters.

Parameter bindings are stored in the `substitution_parameters` folder inside the data generator directory. Each query gets its bindings in a separate file. Every line of a parameter file is a JSON-formatted collection of key-value pairs (name of the parameter and its value). For example, the Query 1 parameter bindings are stored in file `query_1_param.txt`, and one of its lines may look like this:

{"PersonID" : 1, "Name" : "Lei", "PersonURI" : "http://www.ldbc.eu/ldbc_socialnet/1.0/data/pers1"}

Depending on implementation, the SUT may refer to persons either by IDs (relational and graph databases) or URIs (RDF systems), so we provide both values for the Person parameter. Finally, parameters for short reads are taken from those in complex reads and updates.

4.4 Load Definition

LDBC-SNB Test Driver is in charge of the execution of the Interactive Workload. At the beginning of the execution, the Test Driver creates a query mix by assigning to each query instance, a query issue time and a set of parameters taken from the generated substitution parameter set described above.

Query issue times have to be carefully assigned. Although substitution parameters are chosen in such a way that queries of the same type take similar time, not all query types have the same complexity and touch the same

amount of data, which causes them to scale differently for the different scale factors. Therefore, if all query instances, regardless of their type, are issued at the same rate, those more complex queries will dominate the execution's result, making faster query types purposeless. To avoid this situation, each query type is executed at a different rate. The way the execution rate is decided, also depends on the nature of the query: complex read, short read or update.

Update queries' issue times are taken from the update streams generated by the data generator. These are the times where the actual event happened during the simulation of the social network. Complex reads' times are expressed in terms of update operations. For each complex read query type, a frequency value is assigned which specifies the relation between the number of updates performed per complex read. Table 4.1 shows the frequencies assigned to each query type for SF1. The frequencies of the different scale factors can be found in Appendix A.1.

Query Type	freq	Query Type	freq
Query 1	26	Query 8	45
Query 2	37	Query 9	157
Query 3	69	Query 10	30
Query 4	36	Query 11	16
Query 5	57	Query 12	44
Query 6	129	Query 13	19
Query 7	87	Query 14	49

Table 4.1: Frequencies for each query type for SF1.

Finally, short reads are inserted in order to balance the ratio between reads and writes, and to simulate the behavior of a real user of the social network. For each complex read instance, a sequence of short reads is planned. There are two types of short read sequences: Person centric and Message centric. Depending on the type of the complex read, one of them is chosen. Each sequence consists of a set of short reads which are issued in a row. The issue time assigned to each short read in the sequence is determined at run time, and is based on the completion time of the complex read it depends on. The substitution parameters for short reads are taken from the results of previously executed complex reads and short reads. Once a short read sequence is issued (and provided that sufficient substitution parameters exist), there is a probability that another short read sequence is issued. This probability decreases for each new sequence issued. Since the same random number generator seed is used across executions, the workload is deterministic.

The specified frequencies, implicitly define the query ratios between queries of different types, as well as a default target throughput. However the Test Sponsor may specify a different target throughput to test, by "squeezing" together or "stretching" apart the queries of the workload. This is achieved by means of the "Time Compression Ratio" that is multiplied by the frequencies (see Table 4.1). Therefore, different throughputs can be tested while maintaining the relative ratios between the different query types.

5.1 Choke Points

5.1.1 Aggregation Performance

SNBI-1.1/TCPH-1.2: [QOPT] Interesting Orders

This choke-point tests the ability of the query optimizer to exploit the interesting orders induced by some operators. Apart from clustered indexes providing key order, other operators also preserve or even induce tuple orderings. Sort-based operators create new orderings, typically the probe-side of a hash join conserves its order, etc.

SNBI-1.2/TCPH-1.1: [QEXE] High Cardinality group-by performance

This choke-point tests the ability of the execution engine to parallelize group-by's with a large number of groups. Some queries require performing large group-by's. In such a case, if an aggregation produces a significant number of groups, intra query parallelization can be exploited as each thread may make its own partial aggregation. Then, to produce the result, these have to be re-aggregated. In order to avoid this, the tuples entering the aggregation operator may be partitioned by a hash of the grouping key and be sent to the appropriate partition. Each partition would have its own thread so that only that thread would write the aggregation, hence avoiding costly critical sections as well. A high cardinality distinct modifier in a query is a special case of this choke point. It is amenable to the same solution with intra query parallelization and partitioning as the group-by. We further note that scale-out systems have an extra incentive for partitioning since this will distribute the CPU and memory pressure over multiple machines, yielding better platform utilization and scalability.

SNBI-1.3: [QEXE] Complex aggregate performance

This choke-point test the performance of the execution engine to perform complex aggregates. Many databases offer user defined aggregates and more complex aggregation operations than the basic count, sum, max and min, for example string concatenation aggregation operator. These types of aggregates are expected to benefit from the same optimizations as the basic built-in ones, for example partitioning.

SNBI-1.4: [QOPT] Top-k push down

Top-k push down. This choke-point tests the ability of the query optimizer to perform optimizations based on top-k selections. Many times queries demand for returning the top-k elements. Once k results are obtained, extra restrictions in a selection can be added based on the properties of the kth element currently in the top-k, being more restrictive as the query advances, instead of sorting all elements and picking the highest k.

SNBI-1.5/TCPH-1.4: [QEXE] Dependent group-by keys

This choke-point tests the ability of the query optimizer to exclude those functionally dependent group-bys. Sometimes queries require for group-by's on a set of columns and a key, where the value of the key determines the columns. In this situation, the aggregation operator should be able to exclude certain group-by attributes from the key matching.

SNBI-1.6/TCPH-1.3: [QEXE] Low Cardinality group-by performance

This choke-point tests the ability to efficiently perform group by evaluation when only a very limited set of groups is available. This can require special strategies for parallelization, e.g. pre-aggregation when possible. This case also allows using special strategies for grouping like using array lookup if the domain of keys is small.

5.1.2 Join Performance

SNBI-2.1/TCPH-2.3: [QOPT] Rich join order optimization

This choke-point tests the ability of the query optimizer to find optimal join orders. A graph can be traversed in different ways. In the relational model, this is equivalent as different join orders. The execution time of these orders may differ by orders of magnitude. Therefore, finding an efficient join (traversal) order is important, which in general, requires enumeration of all the possibilities. The enumeration is complicated by operators that are not freely re-orderable like semi-, anti-, and outer-joins. Because of this difficulty most join enumeration algorithms do not enumerate all possible plans, and therefore can miss the optimal join order. Therefore, these chokepoint tests the ability of the query optimizer to find optimal join (traversal) orders.

SNBI-2.2/TCPH-2.4: [QOPT] Late projection

This choke-point tests the ability of the query optimizer to delay the projection of unneeded attributes until late in the execution. Queries where certain columns are only needed late in the query. In such a situation, it is better to omit them from initial table scans, as fetching them later by row-id with a separate scan operator, which is joined to the intermediate query result, can save temporal space, and therefore I/O. Late projection does have a trade-off involving locality, since late in the plan the tuples may be in a different order, and scattered I/O in terms of tuples/second is much more expensive than sequential I/O. Late projection specifically makes sense in queries where the late use of these columns happens at a moment where the amount of tuples involved has been considerably reduced; for example after an aggregation with only few unique group-by keys, or a top-k operator.

SNBI-2.3/TCPH-2.1: [QOPT] Join type selection

This choke-point tests the ability of the query optimizer to select the proper join operator type, which implies accurate estimates of cardinalities. Depending on the cardinalities of both sides of a join, a hash or an index index based join operator is more appropriate. This is especially important with column stores, where one usually has an index on everything. Deciding to use a hash join requires a good estimation of cardinalities on both the probe and build sides. In TPC-H, the use of hash join is almost a foregone conclusion in many cases, since an implementation will usually not even define an index on foreign key columns. There is a break even point between index and hash based plans, depending on the cardinality on the probe and build sides

SNBI-2.4/TCPH-2.2: [QOPT] Sparse foreign key joins

This choke-point tests the performance of join operators when the join is sparse. Sometimes joins involve relations where only a small percentage of rows in one of the tables is required to satisfy a join. When tables are larger, typical join methods can be sub-optimal. Partitioning the sparse table, using Hash Clustered indexes or implementing bloom filter tests inside the join are techniques to improve the performance in such situations [3].

5.1.3 Data Access Locality

SNBI-3.1/TCPH-3.3: [QOPT] Detecting correlation

This choke-point tests the ability of the query optimizer to detect data correlations and exploiting them. If a schema rewards creating clustered indexes, the question then is which of the date or data columns to use as key. In fact it should not matter which column is used, as range- propagation between correlated attributes of the same table is relatively easy. One way is through the creation of multi-attribute histograms after detection of attribute correlation. With MinMax indexes, range-predicates on any column can be translated into qualifying tuple position ranges. If an attribute value is correlated with tuple position, this reduces the area to scan roughly equally to predicate selectivity.

SNBI-3.2: [STORAGE] Dimensional clustering

This chokepoint tests suitability of the identifiers assigned to entities by the storage system to better exploit data locality. A data model where each entity has a unique synthetic identifier, e.g. RDF or graph models, has some choice in assigning a value to this identifier. The properties of the entity being identified may affect this, e.g. type (label), other dependent properties, e.g. geographic location, date, position in a hierarchy etc, depending on the application. Such identifier choice may create locality which in turn improves efficiency of compression or index access.

SNBI-3.3: [QEXE] Scattered Index Access patterns

This choke-point tests the performance of indexes when scattered accesses are performed. The efficiency of index lookup is very different depending on the locality of keys coming to the indexed access. Techniques like vectoring non-local index accesses by simply missing the cache in parallel on multiple lookups vectored on the same thread may have high impact. Also detecting absence of locality should turn off any locality dependent optimizations if these are costly when there is no locality. A graph neighborhood traversal is an example of an operation with random access without predictable locality.

5.1.4 Expression Calculation

SNBI-4.1/TPCH-4.2a: [QOPT] Common subexpression elimination

This choke-point tests the ability of the query optimizer to detect common sub-expressions and reuse their results. A basic technique helpful in multiple queries is common subexpression elimination (CSE). CSE should recognize also that average aggregates can be derived afterwards by dividing a SUM by the COUNT when those have been computed.

SNBI-4.2/TCPH-4.2d: [QOPT] Complex boolean expression joins and selections

This choke-point tests the ability of the query optimizer to reorder the execution of boolean expressions to improve the performance. Some boolean expressions are complex, with possibilities for alternative optimal evaluation orders. For instance, the optimizer may reorder conjunctions to test first those conditions with larger selectivity [5].

SNBI-4.3 [QEXE] Low overhead expressions interpretation

This choke-point tests the ability to efficiently evaluate simple expressions on a large number of values. A typical example could be simple arithmetic expressions, mathematical functions like floor and absolute or date functions like extracting a year.

5.1.5 Correlated Sub-queries

SNBI-5.1/TCPH-5.1: [QOPT] Flattening sub-queries

This choke-point tests the ability of the query optimizer to flatten execution plans when there are correlated sub-queries. Many queries have correlated sub-queries and their query plans can be flattened, such that the correlated sub-query is handled using an equi-join, outer-join or anti-join. To execute queries well, systems need to flatten both sub-queries, the first into an equi-join plan, the second into an anti-join plan. Therefore, the execution layer of the database system will benefit from implementing these extended join variants. The ill effects of repetitive tuple-at-a-time sub-query execution can also be mitigated if execution systems by using vectorized, or block-wise query execution, allowing to run sub-queries with thousands of input parameters instead of one. The ability to look up many keys in an index in one API call creates the opportunity to benefit from physical locality, if lookup keys exhibit some clustering.

SNBI-5.2/TCPH-5.3: [QEXE] Overlap between outer and sub-query

This choke-point tests the ability of the execution engine to reuse results when there is an overlap between the outer query and the sub-query. In some queries, the correlated sub-query and the outer query have the same joins and selections. In this case, a non-tree, rather DAG-shaped [6] query plan would allow to execute the common parts just once, providing the intermediate result stream to both the outer query and correlated sub-query, which higher up in the query plan are joined together (using normal query decorrelation rewrites). As such, the benchmark rewards systems where the optimizer can detect this and the execution engine supports an operator that can buffer intermediate results and provide them to multiple parent operators.

SNBI-5.3/TCPH-5.2 ? TODO: [QEXE] Intra-query result reuse

This choke-point tests the ability of the execution engine to reuse sub-query results when two sub-queries are mostly identical. Some queries have almost identical sub-queries, where some of their internal results can be reused in both sides of the execution plan, thus avoiding to repeat computations.

5.1.6   Parallelism and Concurrency

SNBI-6.1/TCPH-6.3: [QEXE] Inter-query result reuse

This choke-point tests the ability of the query execution engine to reuse results from different queries. Sometimes with a high number of streams a significant amount of identical queries emerge in the resulting workload. The reason is that certain parameters, as generated by the workload generator, have only a limited amount of parameters bindings. This weakness opens up the possibility of using a query result cache, to eliminate the repetitive part of the workload. A further opportunity that detects even more overlap is the work on recycling, which does not only cache final query results, but also intermediate query results of a "high worth". Here, worth is a combination of partial-query result size, partial-query evaluation cost, and observed (or estimated) frequency of the partial-query in the workload.

5.1.7   RDF and Graph Specifics

SNBI-7.1: [QOPT] Translation of internal ids to external ones [KILL WITH FIRE]

This choke-point tests the ability of the query optimizer to delay the translation between internal and external entity ids to late in the query. Translate at point of minimum cardinality, e.g. after top k order by RDF and possibly graph models often use a synthetic integer identifier for entities, e.g. URI's . For presentation to the client applications, these identifiers must be translated to their original form, e.g. the URI string that was used when loading the data. This should be done as late as possible, or at the point of minimal cardinality in the plan.

SNBI-7.2: [QOPT] Cardinality estimation of transitive paths

This choke-point tests the ability of the query optimizer to properly estimate the cardinality of intermediate results when executing transitive paths. A transitive path may occur in a "fact table" or a "dimension table" position. A transitive path may cover a tree or a graph, e.g. descendants in a geographical hierarchy vs. graph neighborhood or transitive closure in a many-to-many connected social network. In order to decide proper join order and type, the cardinality of the expansion of the transitive path needs to be correctly estimated. This could for example take the form of executing on a sample of the data in the cost model or of gathering special statistics, e.g. the depth and fan-out of a tree. In the case of hierarchical dimensions, e.g. geographic locations or other hierarchical classifications, detecting the cardinality of the transitive path will allow one to go to a star schema plan with scan of a fact table with a selective hash join. Such a plan will be on the other hand very bad for example if the hash table is much larger than the "fact table" being scanned.

SNBI-7.3: [QEXE] Execution of a transitive step

This choke-point tests the ability of the query execution engine to efficiently execute transitive steps. Graph workloads may have transitive operations, for example finding a shortest path between vertices. This involves repeated execution of a short lookup, often on many values at the same time, while usually having an end condition, e.g. the target vertex being reached or having reached the border of a search going in the opposite direction. For the best efficiency, these operations can be merged or tightly coupled to the index operations themselves. Also parallelization may be possible but may need to deal with a global state, e.g. set of visited vertices. There are many possible tradeoffs between generality and performance

SNBI 7.4: [QEXE] Efficient evaluation of termination criteria for transitive queries

This tests the ability of a system to express termination criteria for transitive queries so that not the whole transitive relation has to be evaluated as well as efficient testing for termination.

SNBI 7.5: [QEXE] Path pattern reuse

This choke-point tests the ability of the execution engine to reuse work across graph traversals [TODO: complete in more detail] For example, when computing paths within a range of distances, it is often possible to incrementally compute longer paths by reusing paths of shorter distances that were already computed

5.2 Query Specifications

number	1
title	Posting summary
	message: Message *creationDate < $date* length year(creationDate)

description	Given a date, find all Messages created before that date. Group them by a 3-level grouping:
	
	1. by year of creation
	2. for each year, group into message types, i.e., Posts or Comments
	3. for each year-type group, split into four groups based on length of their content
	
	    • 0 <= length < 40: `short`
	    • 40 <= length < 80: `one liner`
	    • 80 <= length < 160: `tweet`
	    • 160 <= length: `long`

group	`year`, `message type`, `length group`
parameters	`date` Date
result	`message.year` 32bitInteger `message type` String `length category` String `message count` 32bitInteger `average message length` 32bitInteger `sum message length` 32bitInteger `per messages` 32bitFloat
sort	`year` ↓ `message type` ↑ `size category` ↑
choke points	1.2, 3.2, 4.1

number	2
title	Top tags for country, age, gender, time

description	Select all Messages (Posts & Comments) created between date1-date2 (inclusive) by persons located in country1 or country2. Select the creators (Persons) and the Tags of these Messages. Split these Persons, Tags and Messages into a 5-level grouping: (1) name of country of person, (2) month message was created, (3) gender of person, (4) age group of person, defined as years between person's birthday and end of simulation (2013-01-01), divided by 5, rounded down, (5) name of tag attached to message. Only return groups where number of messages is greater than 100.
group	`countryName, month, gender, ageGroup, tagName`
parameters	`date1` Date `date2` Date `country1` String `country2` String
result	`country.name` String `message.month` 32bitInteger `person.gender` String `ageGroup` 32bitInteger `tag.name` String `messageCount` 64bitInteger
sort	`messageCount` ↓ `tag.name` ↑ `ageGroup` ↑ `person.gender` ↑ `message.month` ↑ `country.name` ↑
limit	100
choke points	1.1, 1.2, 1.4, 2.1, 2.3, 3.1, 3.2

number	3
title	Tag evolution

description	Given a year and a month, find the Tags that were used in Messages during the given month of the given year, and the Tags that were used during the month after the given month of the given year. For both months, compute the count of Messages that used each of the Tags.
parameters	year 32bitInteger month 32bitInteger
result	tag.name String countMonth1 32bitInteger occurrences of the tag during year-month 1 countMonth2 32bitInteger occurrences of the tag during year-month 2 diff 32bitInteger difference between occurrences of this Tag in month 1 and month 2
sort	diff ↓ tag.name ↑
limit	100
choke points	2.4, 3.1, 3.2, 4.1, 4.3, 5.3, 6.1

number	4
title	Popular topics in a country

description	Given a TagClass and a Country, find all the Forums created in the given Country, containing at least one Post with Tags belonging directly to the given TagClass. The location of a Forum is identified by the location of the Forum's moderator. TODO - what do we count, Posts? (szarnyasg)
parameters	tagClass 32bitInteger country 32bitInteger
result	forum.id 64bitInteger forum.title String forum.creationDate DateTime person.id 64bitInteger count 32bitInteger
sort	count ↓ forum.id ↑
limit	20
choke points	1.1, 1.2, 1.4, 2.1, 2.2, 2.4, 3.3

number	5
title	Top posters in a country

description	Find the most popular Forums for a given Country, where the popularity of a Forum is measured by the number of members that Forum has from the given Country. For each member of the 100 most popular Forums, count the number of Posts they made in any of those (most popular) Forums.
group	`person.id`, `person.firstName`, `person.lastName`, `person.creationDate`
parameters	country 32bitInteger
result	person.id 64bitInteger person.firstName String person.lastName String person.creationDate DateTime postCount 32bitInteger
sort	postCount ↓ person.id ↑
limit	100
choke points	1.2, 1.4, 1.5, 2.1, 2.2, 2.3, 2.4, 3.3, 5.3, 6.1

number	6
title	Most active Posters of a given Topic

description	Get Persons who have created a Message (Post or Comment) with a given Tag. Each Person has a score, computed as follows: • Count of Messages with the given Tag (`postCount`). • Count of Likes (`likeCount`) and Comments (`replyCount`) in reply of their Messages with the given Tag. (TODO - transitive or direct? szarnyasg) The sum is weighted as follows: • Messages (`postCount`) are multiplied by 1, • Comments to Messages (`replyCount`) are multiplied by 2, • Likes (`likeCount`) are multiplied by 10.
parameters	`tag` 32bitInteger
result	`person.id` 64bitInteger `replyCount` 32bitInteger `likeCount` 32bitInteger `postCount` 32bitInteger `score` 32bitInteger
sort	`score` ↓ `person.id` ↑
limit	100
choke points	1.2, 2.3

number	7
title	Most authoritative users on a given topic
description	Given a Tag, find all Persons that ever created a Message with the given Tag. For each of these Persons compute their "authority score" as follows: • The "authority score" is the sum of "popularity scores" of the Persons that liked any of that Person's Messages with the given Tag. • A Person's "popularity score" is defined as the total number of likes on all of their Messages.
parameters	`tag` 32bitInteger
result	`person1.id` 64bitInteger `authorityScore` 32bitInteger
sort	`authorityScore` ↓ `person1.id` ↑
limit	100
choke points	1.2, 2.3, 3.2, 3.3, 6.1

number	8
title	Related topics

description	Find all Messages that have a given Tag. Find the related Tags attached to replies of these Messages (TODO - transitive? szarnyasg), but only of those replies that do not have the given Tag. Group the Tags by name, and get the count of replies in each group.
group	`relatedTag.name`
parameters	`tag` 32bitInteger
result	`relatedTag.name` String

`count` 32bitInteger |
| sort | `count` ↓

`relatedTag.name` ↑ |
| limit | 100 |
| choke points | 1.6, 3.3, 5.2 |

number	9
title	Forum with related Tags
description	Given two TagClasses (`tagClass1` & `tagClass2`), find Forums that contain at least one Post with a Tag from `tagClass1` and at least one Post with a Tag from `tagClass2` (direct children not transitive) – this may be the same Post. Consider the Forums with a number of members greater than a given threshold. For every such forum, count the number of Posts that have a Tag from TagClass1 (count1), and the number of posts that have a tag from TagClass2.
parameters	tagClass1 `32bitInteger` tagClass2 `32bitInteger` threshold `32bitInteger`
result	forum.id `64bitInteger` count1 `32bitInteger` Number of Posts with at least one tag belonging to tagClass1 count2 `32bitInteger` Number of Posts with at least one tag belonging to tagClass2
sort	count2 ↓ count1 ↓ forum.id ↑
limit	100
choke points	1.2, 1.4, 2.1, 2.3, 2.4

number	10
title	Central Person for a Tag

description	Given a Tag, find all Persons that are either interested in the Tag, or have written a Message (Post or Comment) with creation date after a given date and that has a given Tag. For each Person, compute the score as the sum of the following two aspects: • 100, if the Person has this tag as their interest, or 0 otherwise • number of messages by this person with the given tag
parameters	`tag` 32bitInteger `date` Date
result	`person.id` 64bitInteger `score` 32bitInteger `friendsScore` 32bitInteger The sum of the score of the Person's friends
sort	`score + friendsScore` ↓ `person.id` ↑
limit	100
choke points	1.2, 2.1, 2.3, 3.2

number	11
title	Unrelated replies

description	Find those Persons of a given country that replied to any Message, such that the reply does not have any Tag in common with the Message [are transitive replies considered or not? - SzG]. Consider only those replies not containing any word from a given blacklist. For each Person and valid reply, retrieve the Tags associated with the reply, and retrieve the number of likes on the reply.
group	`person.id`, `tag.name`
parameters	`country` 32bitInteger `blacklist` String[]
result	`person.id` 64bitInteger `tag.name` String `countLikes` 32bitInteger The count of Likes to replies with that Tag. `countReplies` 32bitInteger The count of replies with that Tag.
sort	`countLikes` ↓ `person.id` ↑ `tag.name` ↑
limit	100
choke points	1.1, 2.1, 2.2, 2.3, 3.1, 3.2, 6.1

number	12
title	Trending Posts
description	Find all Messages created after a given date, that received more than a given number of likes.
parameters	`creationDate` Date `likeThreshold` 32bitInteger
result	`message.id` 64bitInteger `message.creationDate` DateTime `creator.firstName` String The first name of the post's creator `creator.lastName` String The last name of the post's creator `likeCount` 32bitInteger The number of Likes the Post received
sort	`likeCount` ↓ `message.id` ↑
limit	100
choke points	1.2, 2.2, 3.1, 6.1

number	13
title	Popular Tags per month in a country
description	Find all Messages in a given Country, as well as their Tags. For each group, find the 5 most popular Tags, where popularity is the number of Messages (from within the same group) where the Tag appears.
group	`year, month`
parameters	`country` String
result	`year` 32bitInteger year(message.creationDate) `month` 32bitInteger month(message.creationDate) `popularTags` TagPairs (tag.name - String, popularity - 32bitInteger), sorted descending by popularity, then ascending by tag name
sort	`year` ↓ `month` ↑
limit	100
choke points	1.2, 2.2, 2.3, 3.2, 6.1

number	14
title	Top thread initiators

description	For each person, count the number Posts they created in the time interval (`begin, end`), and the number of messages in each of their (transitive) reply trees. When calculating message counts only consider messages created within the given time interval. Return each person, number of Posts they created, and the count of all messages that appeared in the reply trees (including Post at tree root) they created.
parameters	`begin` Date `end` Date
result	`person.id` 64bitInteger `person.firstName` String `person.lastName` String `threadCount` 32bitInteger The number of threads initiated by that Person `messageCount` 32bitInteger The number of messages created in all the threads this Person initiated
sort	`messageCount` ↓ `person.id` ↑
limit	100
choke points	1.2, 2.2, 2.3, 3.2, 7.2, 7.3, 7.4
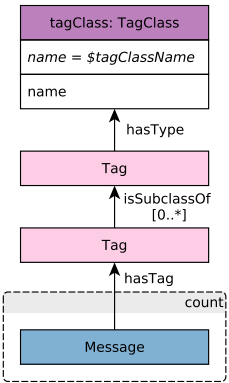
number	15
title	Social normals

description	Given a country, find all Persons of the country whose number of friends in the given country equals the (floor of) average number of friends that Persons of the given country have in the given country.
parameters	country 32bitInteger
result	person.id 64bitInteger count 32bitInteger
sort	person.id ↑
limit	100
choke points	1.2, 2.3, 3.2, 3.3, 5.3, 6.1

number	16
title	Experts in social circle

description	Given a Person, find all other Persons that live in a given country and are connected to given person by a transitive path with length in range `[min, max]` through the knows relation. [DISCUSS: edge isomorphism semantics] For each of these Persons, retrieve all of their Messages (Posts & Comments) that contain at least one Tag belonging to a given TagClass (direct relation not transitive). For each Message, also retrieve its Tags. TODO [szarnyasg]: what is postCount?
group	`tag.name`, `person.id`
parameters	`personId` 64bitInteger `country` String `tagClass` String `minPathDistance` 32bitInteger `maxPathDistance` 32bitInteger
result	`person.id` 64bitInteger `tag.name` String `messageCount` 32bitInteger number of Messages created by that Person containing that Tag
sort	`messageCount` ↓ `tag.name` ↑ `person.id` ↑
limit	100
choke points	1.2, 1.4, 2.3, 2.4, 3.3, 7.2, 7.3

number	17
title	Friend triangles

description	For a given country, count all the distinct triples of persons such that `a` is friend of `b`, `b` is friend of `c`, and `c` is friend of `a`. Distinct means that given a triple *t1* in the result set *R* of all qualified triples, there is not a triple *t2* in *R* such that |*t1* U *b*| = 3.
parameters	country String
result	count 32bitInteger
choke points	1.1, 2.3

number	18
title	How many persons have a given number of posts

description	For each Person, count the number of Messages (Posts & Comments) they made. Only consider messages with: • length below the `lengthThreshold` • creationDate after `creationDate` (TODO - is after exclusive or inclusive, does it allow equality?) • any of the given `languages` (TODO - only Posts have a language, messages do not)
group	`messageCount`
parameters	`creationDate` Date `lengthThreshold` TODO (32bitInteger?)
result	`messageCount` 32bitInteger number of messages created `personCount` 32bitInteger the number of Persons with 'messageCount' messages
sort	`personCount` ↓
choke points	1.1, 1.2, 1.6, 3.2, 4.2, 4.3

number	19
title	Stranger's interaction
description	For all the Persons born after a certain date, find all the strangers they interacted with, where strangers are Persons that do not Know each other. There is no restriction on the date that strangers were born. Consider only strangers that are members of Forums tagged with tagClass1 (direct children not transitive) AND members of Forums tagged with tagClass2 (direct children not transitive). It does not matter if these Tags are attached to the same Forum, or different Forums. We define interaction as follows: a Person replies to a Message (Post or Comment) by another Person. For each Person, count the number of strangers they interacted (directed) with and total number of times they interacted (directed) with them.
parameters	date Date tagClass1 32bitInteger tagClass2 32bitInteger
result	person.id 64bitInteger strangersCount 32bitInteger interactionCount 32bitInteger
sort	interactionCount ↓ person.id ↑
limit	100
choke points	1.1, 1.4, 2.1, 2.3, 2.4, 3.3, 5.1, 7.3, 7.4

number	20
title	High-level topics

UNWIND $tagClassNames[] -> $tagClassName

description	For all given TagClasses, count number of Messages that have a Tag that belongs to that TagClass or any of its children (all descendants through a transitive relation).
parameters	`tagClasses` `String[]`
result	`tagClass.name` `String` the TagClass of the root `postCount` `32bitInteger`
sort	`postCount` ↓ `tagClass.name` ↑
limit	100
choke points	1.6, 2.1, 6.1

number	21
title	Zombies in a country

description	Find zombies within the given country, and return their zombie scores. A zombie is a Person created before the given `endDate` and that has created between `[0, 1)` Messages per month, on average, during the time range between profile creation date and the given `endDate`. The number of months spans the time range from creation date of the profile to the `endDate` and also includes partial months.
parameters	`country` String `endDate` Date
result	`person.id` 64bitInteger `zombieLikeCount` 32bitInteger `realLikeCount` 32bitInteger `zombieScore` 32bitFloat the ratio between the number of likes received from zombies ('zombieLikeCount') and the total number of likes received on that person's messages ('realLikeCount') – only count likes received from profiles that were created before the given 'endDate'.
sort	`zombieScore` ↓ `person.id` ↑
limit	100
choke points	1.2, 2.1, 2.3, 2.4, 3.2, 3.3, 5.1, 5.3

number	22
title	International dialog

description	Consider all pairs of people (`p1`, `p2`) such that one is located in a city of Country `countryX` and the other is located in a city of Country `countryY`. For each city of Country `countryX`, return the highest scoring pair. The score of a pair is defined as the sum of the scores of the following kinds of interaction: • `p1` has created a reply Comment to at least one Comment or Post by `p2`: `Score = 4` • `p1` has created at least one Post or Comment that `p2` has created a reply Comment to: `Score = 1` • `p1` and `p2` Know each other: `Score = 15` • `p1` liked at least one Post or Comment by `p2`: `Score = 10` • `p1` has created at least one Post or Comment that was liked by `p2`: `Score = 1` I.e., the maximum score a pair can obtain is: `4 + 1 + 15 + 10 + 1 = 31`
parameters	`countryX` String `countryY` String
result	`p1.id` 64bitInteger `p2.id` 64bitInteger `cityX.name` String `score` 32bitInteger
sort	`score` ↓ `p1.id` ↑ `p2.id` ↑
choke points	1.4, 1.6, 2.1, 3.1, 3.3, 5.1, 5.2, 5.3

number	23
title	Holiday destinations

description	Count the messages all residents of Country `country` have written abroad grouped by month and Country. A Message was written abroad if the Country the Message was written in is different than the Country of the Person it was written by.
parameters	`country` String
result	`messageCount` 32bitInteger The number of messages in each group `country.name` String The name of the destination country `month` 32bitInteger
sort	`messageCount` ↓ `country.name` ↑ `month` ↑
limit	100
choke points	1.6, 2.3, 2.4, 3.3, 4.3

number	24
title	Messages by Topic and Continent

description	Find all Messages tagged with a Tag from the given TagClass (non-transitive). Count all messages and their likes grouped by continent, year, and month. (TODO - do we group the Messages or the Persons who liked the Messages by continent? I think the former one - szarnyasg)
group	`year, month, continent.name`
parameters	`tagClass` String
result	`messageCount` 32bitInteger `likeCount` 32bitInteger `year` 32bitInteger year of the Message's creationDate `month` 32bitInteger month of the Message's creationDate `continent.name` String
sort	`year` ↑ `month` ↑ `continent.name` ↓
limit	100
choke points	1.6, 2.1, 2.3, 2.4, 3.2, 4.3

number	25
title	Weighted paths

description	Given two Persons, find all (unweighted) shortest paths between these two Persons, in the subgraph induced by the Knows relationship.
	Then, for each path calculate a weight. The nodes in the path are Persons, and the weight of a path is the sum of weights between every pair of consecutive Person nodes in the path.
	The weight for a pair of Persons is calculated such that every reply (by one of the Persons) to a Post (by the other Person) contributes 1.0, and every reply (by ones of the Persons) to a Comment (by the other Person) contributes 0.5. Only consider messages that were created in a forum that was created within the timeframe `[startDate, endDate]`.
	Return all the paths with shortest length, and their weights.
parameters	`person1Id` 64bitInteger
	`person2Id` 64bitInteger
	`startDate` Date
	`endDate` Date
result	`person.id` 64bitInteger Identifiers representing an ordered sequence of the Persons in the path weight
sort	`weight` ↓ The order of paths with the same weight is unspecified
choke points	1.2, 2.1, 2.2, 2.4, 3.3, 5.1, 5.3, 7.2, 7.3

6 AUDITING RULES

This chapter describes the rules to audit benchmark runs, that is, what techniques are allowed and what are not, what must be provided to the auditor and guidelines for the auditors to perform the audit.

6.1 Preparation

The first step when doing an audit is to determine the versions of the following items that will be used for the benchmark:

- The benchmark specification

- The data generator

- The driver

These must be reported in the full disclosure report to guarantee that the benchmark run can be reproduced exactly in the future. Similarly, the test sponsor will inform the auditor the scale factor to test. Finally, a clean test system with enough space to store the scale factor must be provided, including the update streams and substitution parameters.

6.1.1 Collect System Details

The next step is to collect the technical and pricing details of the system under test. This includes the following items:

- Common name of the system, e.g. Dell PowerEdge xxxx.

- Type and number of CPUs, cores/threads per CPU, clock frequency and cache hierarchy characteristics (levels, size per level, etc.).

- The amount of the system's memory, type and frequency.

- The disk controller or motherboard type if disk controller is on the motherboard.

- For each distinct type of secondary storage device, the number and characteristics of the device.

- The number and type of network controllers.

- The number and type of network switches. Wiring must be disclosed.

- Date of availability of the system.

Only the network switches and interfaces that participate in the run need to be reported. If the benchmark execution is entirely contained on a single machine, no network need be reported. The price of the hardware in question must be disclosed and should reflect the single quantity list price that any buyer could expect when purchasing one system with the given specification. The price may be either an item by item price or a package price if the system is sold as a package

Besides hardware characteristics, also software details must be collected:

- The DBMS and operating system name and versions.

- Installation and configuration information of both the DBMS and operating system, which must be provided by the test sponsor.

- Price of the software license used, which can be tied to the number of concurrent users or size of data.

- Date of availability of the software.

Also, the test sponsor must provide all the source code relevant to the benchmark.

　　　　　　　　　　　　　LDBC Social Network Benchmark (SNB) – 0.3.0

6.1.2 Setup the Benchmark Environment

Once all the information has been collected, the auditor will setup the environment to perform the benchmark run. This setup includes configuring the following items:

- Setup the LDBC Data generator in the test machine if datasets are not available from a trusted source.

- Setup the LDBC driver with the connectors provided by the test sponsor. The test sponsor must provide the configuration parameters to configure the driver (tcr, number of threads, etc.).

  The `ldbc.snb.interactive.update_interleave` driver parameter must come from the `updateStream.properties` file, which is created by the data generator. That parameter should never be set manually. Also, make sure that the `-rl/--results_log` is enabled. Make sure that all operations are enabled and the frequencies are those for the selected scale factor. These can found in Appendix A.1. If the driver will be executed on a separate machine, gather the characteristics of that machine in the same way as specified above.

6.1.3 Load Data

The test sponsor must provide all the necessary documentation and scripts to load the dataset into the database to test. The system under test must support the different data types needed by the benchmark for each of the attributes at their specified precision. No data can be filtered out, everything must be loaded. The test sponsor must provide a tool to perform arbitrary checks of the data or a shell to issue queries in a declarative language if the system supports it. The auditor will measure the time to load the data, which will be disclosed.

6.2 Running the Benchmark

Running the benchmark consists of three separate parts: (1) validating the query implementations, (2) warming the database and (3) performing the benchmark run. The queries are validated by means of the official validation datasets provided by LDBC consortium in their official software repositories. The auditor must load the provided dataset and run the driver in validation mode, which will test that the queries provide the official results.

The warmup can be performed either using the LDBC driver or externally, and the way it is performed must be disclosed.

A valid benchmark run must last at least 2 hours of simulation time (datagen time). Also, in order to be valid, a benchmark run needs to meet the following requirements. The `results_log.csv` file contains the *actual_start_time* and the *scheduled_start_time* of each of the issued queries. In order to have a valid run, 95% of the queries must meet the following condition:

$$actual\_start\_time - scheduled\_start\_time < 1 \text{ second}$$

If the execution of the benchmark is valid, the auditor must retrieve all the files from directory specified by `-rd/--results_dir` which includes configuration settings used, results log, results summary, which will be disclosed.

6.3 Recovery

Once an official run has been validated, the recovery capabilities of the system must be tested. The system and the driver must be configured in the same way as in during the benchmark execution. After a warmup period, an execution of the benchmark will be performed under the same terms as in the previous measured run.

At an arbitrary point close to 2 hours of simulation execution time, the machine will be disconnected. Then, the auditor will restart the database system and will check that the last commited update (in the driver log file) is actually in the database. The auditor will measure the time taken by the system to recover from the failure. Also, all the information about how durability is ensured must be disclosed. If checkpoints are used, these must be performed with a period of 10 minutes at most.

6.4 Serializability

Optionally, the test sponsor can execute update queries atomically. The auditor will verify that serializability is guaranteed.

References

[1] R. Angles, P. A. Boncz, J. Larriba-Pey, I. Fundulaki, T. Neumann, O. Erling, P. Neubauer, N. Martínez-Bazan, V. Kotsev, and I. Toma. The Linked Data Benchmark Council: a graph and RDF industry benchmarking effort. *SIGMOD Record*, 43(1):27–31, 2014.

[2] O. Erling, A. Averbuch, J. Larriba-Pey, H. Chafi, A. Gubichev, A. Prat-Pérez, M. Pham, and P. A. Boncz. The LDBC Social Network Benchmark: Interactive workload. In *SIGMOD*, pages 619–630, 2015.

[3] G. Graefe. Query evaluation techniques for large databases. *ACM Comput. Surv.*, 25(2):73–170, 1993.

[4] A. Iosup, T. Hegeman, W. L. Ngai, S. Heldens, A. Prat-Pérez, T. Manhardt, H. Chafi, M. Capota, N. Sundaram, M. J. Anderson, I. G. Tanase, Y. Xia, L. Nai, and P. A. Boncz. LDBC Graphalytics: A benchmark for large-scale graph analysis on parallel and distributed platforms. *PVLDB*, 9(13):1317–1328, 2016.

[5] G. Moerkotte. Small materialized aggregates: A light weight index structure for data warehousing. In *PVLDB*, pages 476–487, 1998.

[6] T. Neumann and G. Moerkotte. A framework for reasoning about share equivalence and its integration into a plan generator. In *BTW*, pages 7–26, 2009.

A SCALE FACTOR STATISTICS

A.1 Scale Factor Statistics

Query Type	SF1	SF3	SF10	SF30	SF100	SF300	SF1000
Query 1	26	26	26	26	26	26	26
Query 2	37	37	37	37	37	37	37
Query 3	69	79	92	106	123	142	165
Query 4	36	36	36	36	36	36	36
Query 5	57	61	66	72	78	84	91
Query 6	129	172	236	316	434	580	796
Query 7	87	72	54	48	38	32	25
Query 8	45	27	15	9	5	3	1
Query 9	157	209	287	384	527	705	967
Query 10	30	32	35	37	40	44	47
Query 11	16	17	19	20	22	24	26
Query 12	44	44	44	44	44	44	44
Query 13	19	19	19	19	19	19	19
Query 14	49	49	49	49	49	49	49

Table A.1: Frequencies for each query and SF.